

Build a RESTful API that manages Trase Agents running tasks in production.

At Trase, an Agent is a runnable capability—code, tools, and prompts—designed to complete a class of jobs safely and reliably. A Task is a concrete job instance describing what to do, who can do it, and its lifecycle state. The Agents & Tasks API manages both and coordinates which agents can run which tasks.

- Using any programming language and framework of your choice, create a simple RESTful API that exposes two endpoints: `/agents` and `/tasks`.
- The `/agents` endpoint should support the following operations:
 - o `GET /agents`: Return a list of all agents in JSON format, with each user having an `id`, a `name`, and a `description` attribute.
 - o `POST /agents`: Create a new agent with the given `name` and `description` in the request body, and return the created agent in JSON format, with a unique `id` assigned by the server.
 - o `GET /agents/{id}`: Return the agent with the given `id` in JSON format, or a 404 error if the agent does not exist.
 - o `PUT /agents/{id}`: Update the agent with the given `id` with the new `name` and `description` in the request body, and return the updated agent in JSON format, or a 404 error if the agent does not exist.
 - o `DELETE /agents/{id}`: Delete the agent with the given `id`, and return a 204 status code, or a 404 error if the agent does not exist.
- The `/tasks` endpoint should support the following operations:
 - o `GET /task`: Return a list of all tasks in JSON format, with each post having an `id`, a `title`, a `description`, and a `supported_agent_id` attribute, which references the `id` of the agent who can complete the task. **Bonus:** Make it so tasks can support multiple agents running them.
 - o `POST /tasks`: Create a new task with the given `title`, `description`, and `supported_agent_id` in the request body, and return the created task in JSON format, with a unique `id` assigned by the server. If the `supported_agent_id` does not match any existing agent, return a 400 error.
 - o `GET /tasks/{id}`: Return the task with the given `id` in JSON format, or a 404 error if the task does not exist.
 - o `PUT /tasks/{id}`: Update the task with the given `id` with the new `title`, `description`, and `supported_agent_id` in the request body, and return the updated task in JSON format, or a 404 error if the task does not exist. If the `supported_agent_id` does not match any existing agent, return a 400 error.

- o DELETE /tasks/{id}: Delete the task with the given id, and return a 204 status code, or a 404 error if the task does not exist.
- You can use any data storage method of your choice (in-memory, a file system, or a relational or non-relational database)
- You should provide clear and concise documentation for your API, including the expected request and response formats, the possible error codes and messages, and any assumptions or limitations you made.
- You should also write unit tests and integration tests for your API, using any testing framework of your choice, and provide instructions on how to run them.

Bonus (optional):

- Support soft deletion of agents and tasks
- Create endpoint(s) and data model(s) to enable users to run a task with a particular agent. The user should be able to start a task and get a list of currently running tasks.
- Include a README file explaining how to set up and run the API, and test everything
- Add rate limiting to your API